

```

In [2]: ### import

### basic
import matplotlib.pyplot as plt
import numpy as np
import scipy
import pandas as pd
import math

import seaborn as sns; sns.set()
##matplotlib inline

##sklearn Learners
from sklearn.linear_model import LinearRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.neighbors import KNeighborsRegressor

##sklearn metrics
from sklearn.metrics import accuracy_score
from sklearn.metrics import mean_squared_error
from sklearn.metrics import confusion_matrix

##sklearn model selection
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import validation_curve
from sklearn.model_selection import GridSearchCV

```

Question1

```

In [3]: cd = pd.read_csv("http://www.rob-mcculloch.org/data/susedcars.csv")
cd = cd[['price', 'mileage', 'year']]
cd['price'] = cd['price']/1000
cd['mileage'] = cd['mileage']/1000

```

```

In [46]: X = cd.iloc[:,[1]].to_numpy()
y = cd['price'].to_numpy()
X.shape

```

```

Out[46]: (1000, 1)

```

```

In [6]: rng = np.random.RandomState(34)
Xtrain, Xtest, ytrain, ytest = train_test_split(X,y,random_state=rng, test_size=.2)

kvec = np.arange(348) + 2 #values of k to try
ormsev = np.zeros(len(kvec)) # storage for oos rsmse
irmsev = np.zeros(len(kvec)) # storage for in-sample rsmse

for i in range(len(kvec)):

```

```

# print(i)
tmod = KNeighborsRegressor(n_neighbors=kvec[i])
tmod.fit(Xtrain,ytrain)
yhat = tmod.predict(Xtest)
ormsev[i] = math.sqrt(mean_squared_error(ytest,yhat))
yhat = tmod.predict(Xtrain)
irmsev[i] = math.sqrt(mean_squared_error(ytrain,yhat))


```

```

In [8]: # plot rmse vs k
plt.scatter(kvec,ormsev,c='blue')
plt.plot(kvec,irmsev,c='red')

```

Out[8]: [




```

In [9]: # plot rmse vs model complexity
mcmp = np.log(1/kvec) #model complexity
plt.scatter(mcmp,ormsev,c='blue')
plt.plot(mcmp,irmsev,c='red')
plt.xlabel('model complexity = log(1/k)',size='x-large')
plt.ylabel('rmse',size='x-large')
plt.title('rmse: blue: out of sample, red: in sample',size='x-large')

```

Out[9]: Text(0.5, 1.0, 'rmse: blue: out of sample, red: in sample')




Based on eye-ball method we can choose a k around $k=100$

```
In [25]: tmod = KNeighborsRegressor(n_neighbors=100)
tmod.fit(Xtrain,ytrain)
yhateye = tmod.predict(Xtest)
ormsev = math.sqrt(mean_squared_error(ytest,yhateye))
print(ormsev)
plt.scatter(yhateye,ytest)
plt.plot(ytest,ytest,c='red') #add the line
plt.xlabel('yhateye'); plt.ylabel('y')
```

8.548553226299584

Out[25]: Text(0, 0.5, 'y')



```
In [19]: tempmod = KNeighborsRegressor(n_neighbors=40) #knn with k=40

## rmse from cross validation
cvres = cross_val_score(tempmod,X,y,cv=5,scoring='neg_mean_squared_error') #cross val w

# tranform to rmse
rmse = math.sqrt(np.mean(-cvres))
print('the rmse for k=40 based on 5-fold is:', rmse)

## do it again but shuffle the data
np.random.seed(34)
indices = np.random.choice(X.shape[0],X.shape[0],replace=False)
ys = y[indices]
Xs = X[indices,:]
cvres = cross_val_score(tempmod,Xs,ys,cv=5,scoring='neg_mean_squared_error')
rmse = math.sqrt(np.mean(-cvres))
print('the rmse for k=40 based on 5-fold is:', rmse)
```

```
the rmse for k=40 based on 5-fold is: 9.400370251995131
the rmse for k=40 based on 5-fold is: 9.347582008064498
```


```
In [21]: model = KNeighborsRegressor() # create the knn model

# do cv at every value of k in kvec
trainS, testS = validation_curve(model,X,y,'n_neighbors',kvec,cv=5,scoring='neg_mean_sq

# transform neg_mean_squared_error to rmse
trrmse = np.sqrt(-trainS.mean(axis=1))
termse = np.sqrt(-testS.mean(axis=1))


#plot in and out of sample rmse
plt.scatter(mcmp,termse)
plt.plot(mcmp,trrmse,c='red')
```

```
Out[21]: [<matplotlib.lines.Line2D at 0x1805d917f40>]
```



```
In [22]: plt.scatter(kvec,termse)
plt.plot(kvec,termse,c='red')
```

```
Out[22]: [matplotlib.lines.Line2D at 0x1805d955e20]
```




Based on 5-fold cross validation I think $K=50$ is a good value

```
In [24]: tmod2 = KNeighborsRegressor(n_neighbors=50)
tmod2.fit(Xtrain,ytrain)
yhat = tmod2.predict(Xtest)
ormsev = math.sqrt(mean_squared_error(ytest,yhat))
print(ormsev)
plt.scatter(yhat,ytest)
plt.plot(ytest,ytest,c='red') #add the Line
plt.xlabel('yhtest'); plt.ylabel('y')
```

8.315365111986125

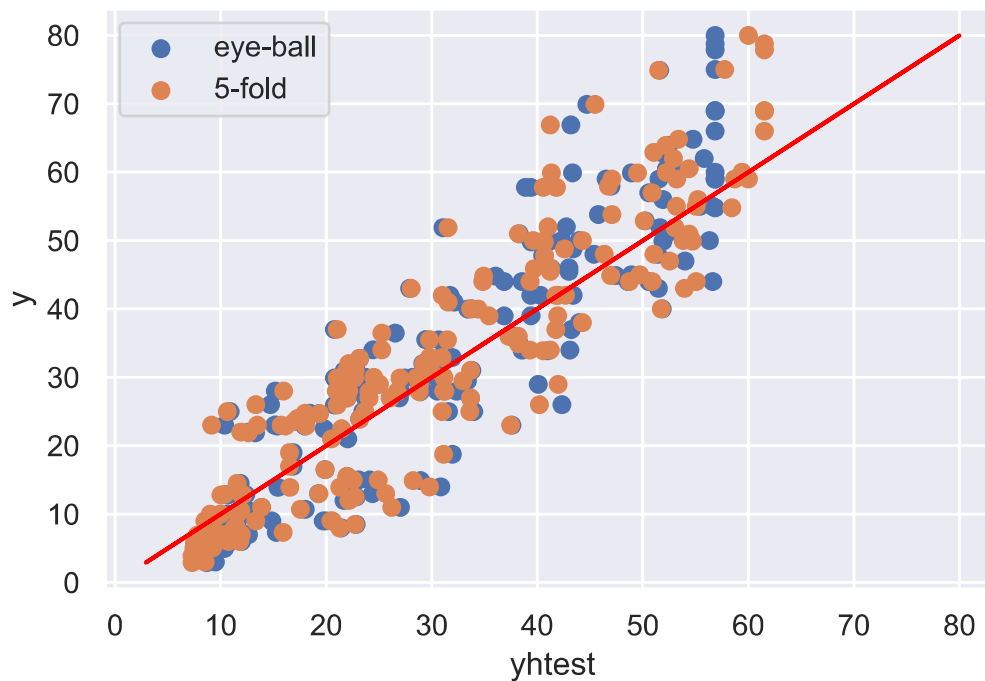
Out[24]: Text(0, 0.5, 'y')



error based on k from 5-fold cross validation is lower than eye-ball method

```
In [29]: plt.scatter(yhateye,ytest,label='eye-ball')
plt.scatter(yhat,ytest,label='5-fold')
plt.plot(ytest,ytest,c='red') #add the Line
plt.xlabel('y_test'); plt.ylabel('y')
plt.legend()
```

Out[29]: <matplotlib.legend.Legend at 0x1805d646e50>



```
In [47]: tmod2 = KNeighborsRegressor(n_neighbors=50)
tmod2.fit(X,y)
xinput=np.array([[100000/1000]])
yhat100 = tmod2.predict(xinput)
print(yhat100)
```

[17.76114]

Question2

```
In [49]: X = cd.iloc[:,[1,2]].to_numpy()
y = cd['price'].to_numpy()
X.shape
```

Out[49]: (1000, 2)


```
In [50]: model = KNeighborsRegressor() # create the knn model

# do cv at every value of k in kvec
trainS, testS = validation_curve(model,X,y,'n_neighbors',kvec,cv=5,scoring='neg_mean_sq

# transform neg_mean_squared_error to rmse
trrmse = np.sqrt(-trainS.mean(axis=1))
termse = np.sqrt(-testS.mean(axis=1))


#plot in and out of sample rmse
plt.scatter(mcmp,termse)
plt.plot(mcmp,trrmse,c='red')
```

Out[50]: [matplotlib.lines.Line2D at 0x1805edb0430]



```
In [32]: plt.scatter(kvec,termse)
plt.plot(kvec,trrmse,c='red')
```

Out[32]: [




```
In [52]: rng = np.random.RandomState(34)
Xtrain, Xtest, ytrain, ytest = train_test_split(X,y,random_state=rng, test_size=.2)
```

```
In [53]: tmod3 = KNeighborsRegressor(n_neighbors=40)
tmod3.fit(Xtrain,ytrain)
yhat = tmod3.predict(Xtest)
ormsev = math.sqrt(mean_squared_error(ytest,yhat))
print(ormsev)
plt.scatter(yhat,ytest)
plt.plot(ytest,ytest,c='red') #add the Line
plt.xlabel('yhtest'); plt.ylabel('y')
```

6.958091773499892

Out[53]: Text(0, 0.5, 'y')



Adding year decreased the error clearly

```
In [54]: tmod3 = KNeighborsRegressor(n_neighbors=40)
tmod3.fit(X,y)
xinput=np.array([[75.0, 2008.]])
yhat75 = tmod3.predict(xinput)
print(yhat75)
```


[29.4746]

Question3

```
In [55]: tmod4 = KNeighborsRegressor(n_neighbors=40, weights='distance')
tmod4.fit(Xtrain,ytrain)
yhat = tmod4.predict(Xtest)
ormsev = math.sqrt(mean_squared_error(ytest,yhat))
print(ormsev)
plt.scatter(yhat,ytest)
plt.plot(ytest,ytest,c='red') #add the Line
plt.xlabel('yhtest'); plt.ylabel('y')
```

6.552581486719303

Out[55]: Text(0, 0.5, 'y')



The new weighting option helps to reduce the error, So it works better

```
In [56]: tmod4 = KNeighborsRegressor(n_neighbors=40, weights='distance')
tmod4.fit(X,y)
xinput=np.array([[75.0, 2008.]])
yhat75 = tmod4.predict(xinput)
print(yhat75)
```

```
[30.52997871]
```